

# Cross-Modal Attention Fusion for Anterior Segment Segmentation Using Complementary Image Representations

Caleb Alfaro  
Computing Engineering  
Instituto Tecnológico de Costa Rica  
calebam2112@gmail.com

Desireé Avilés  
Computing Engineering  
Instituto Tecnológico de Costa Rica  
desiree.aviles.a@gmail.com

Danilo Duque  
Computing Engineering  
Instituto Tecnológico de Costa Rica  
daniloduquedg@gmail.com

**Abstract**—Accurate segmentation of anterior segment structures — including the cornea, pupil, and lens — is clinically relevant for cataract assessment and surgical planning. Unimodal approaches that operate solely on RGB images are sensitive to lighting artifacts and the low color contrast caused by lens opacification, limiting their reliability in real-world ophthalmic settings. We propose a dual-encoder architecture that fuses complementary representations of the same image to address this limitation. One encoder processes the original RGB image to capture photometric and textural features, while a parallel encoder processes a structural edge map derived via Canny filtering to capture geometric boundary information. A cross-attention module at the bottleneck and at encoder stage 4 allows the RGB encoder to query the edge encoder’s feature space, enriching the primary representation with explicit structural cues before decoding. Although both inputs originate from the same image, they occupy heterogeneous feature spaces — photometric and geometric — qualifying the approach as multimodal fusion in the architectural sense. We evaluate the proposed model on the Cataract-Seg dataset against unimodal U-Net baselines and an early-fusion concatenation baseline, reporting Intersection over Union (IoU) and Dice coefficient as primary metrics. The proposed cross-attention fusion strategy achieves competitive performance with strong RGB baselines, matching the best result (Early Fusion, IoU 0.9532) at 0.9530 IoU, while providing a more principled and interpretable fusion mechanism. The near-ceiling performance across RGB-based models indicates dataset saturation; the gated residual design ensures the model does not degrade when edge maps are noisy, making it a practical choice for resource-constrained clinical deployments relative to large foundation model adaptations.

**Index Terms**—cataract segmentation, cross-modal attention, anterior segment, deep learning, multimodal fusion

## I. INTRODUCTION

Cataract disease affects over 100 million people worldwide and remains the leading cause of preventable blindness [1]. Accurately segmenting anterior segment structures, the cornea, pupil, and lens, from clinical images is essential for automated grading, surgical planning, and intraoperative guidance. However, models that process only the RGB image struggle with two problems: the low color contrast caused by lens opacification (the lens becomes cloudy, making its boundary hard to see) and the illumination artifacts common in

slit-lamp photography. These issues reduce reliability precisely in the cases that matter most clinically.

A natural solution is to give the model a second, complementary view of the same image. A *Canny edge map* provides exactly that: it detects boundaries by locating pixels where brightness changes sharply. These boundaries correspond to anatomical edges (cornea boundary, lens capsule) that remain visible even when color contrast is lost. The RGB image carries color and texture information, while the edge map carries boundary information. Although both come from a single photograph, they represent fundamentally different types of information, which makes their fusion a *multimodal* problem in the architectural sense.

Multimodal fusion architectures that combine complementary representations have proven effective in other medical domains. In brain tumor segmentation, cross-attention between different MRI modalities (T1, T2, FLAIR) consistently outperforms both single-modality and early-concatenation baselines [2]. Dual-encoder designs pairing convolutional networks with transformers have shown similar gains [3]. However, these methods assume that multiple modalities come from separate sensors. In ophthalmic imaging, obtaining additional modalities (e.g., OCT alongside a slit-lamp photograph) is not always feasible, and no prior work has explored creating a second representation from a single cataract image.

We propose a dual-encoder U-Net that processes the RGB image through one encoder and its Canny edge map through a second, independent encoder. Cross-attention modules at two stages let the RGB features “query” the edge features for boundary information, enriching the RGB representation with explicit structural cues. A learnable scalar gate controls how much edge information flows in, it starts turned off (initialized so that  $\sigma(\theta) \approx 0.02$ ) so the model initially behaves like a standard RGB U-Net, and gradually learns to rely on edge cues only where they help. Edge encoder skip connections further propagate boundary information through every level of the decoder.

We evaluate on the Cataract-Seg dataset [4] against three baselines: a standard RGB U-Net, an edge-only U-Net, and an early-fusion U-Net that concatenates both inputs at the

first layer. The proposed model achieves an IoU of 0.9530, matching the strongest baseline, while providing a more principled and interpretable fusion mechanism than simple input concatenation.

## II. RELATED WORK

### A. Anterior Segment and Cataract Segmentation

Encoder-decoder architectures have become the standard approach for medical image segmentation since Ronneberger et al. introduced U-Net [5]. U-Net uses a downsampling path followed by a symmetric upsampling path, which enables precise localization even with very few annotated images. This architecture has been widely adopted in ophthalmic segmentation because it preserves fine spatial detail through skip connections.

Surgical scene understanding in cataract procedures presents additional challenges due to reflections, occlusions from instruments, and changes in illumination during surgery. The CaDIS dataset [6] addressed this gap by providing annotated frames of cataract surgery videos covering anatomical structures and surgical instruments, establishing a benchmark for multi-class segmentation in this domain.

### B. Multimodal Fusion in Medical Imaging

The transformer architecture introduced by Vaswani et al. [7] established attention mechanisms as a way to model relationships between different parts of a sequence. Its key insight, that queries from one representation can attend to keys and values from another, directly motivates the cross-modal fusion strategy we adopt at the bottleneck.

In the medical imaging domain, Zhang et al. proposed mmFormer [2], a multimodal medical transformer for brain tumor segmentation that can handle cases where some MRI modalities are missing. mmFormer uses separate encoders for each modality and connects them with attention mechanisms, showing that cross-modal attention outperforms simply concatenating all modalities at the input, especially when different modalities capture very different types of information.

Li et al. proposed DECTNet [3], a dual-encoder architecture combining a convolutional branch with a Swin Transformer branch, fusing their complementary representations via a feature fusion decoder. This confirms that running two parallel encoders and merging their outputs is an effective strategy for medical segmentation. In contrast to DECTNet, which uses two different network architectures processing the same input, our approach uses two identical encoders processing structurally different representations (color and geometry) of the same image.

The multi-modal ophthalmic AI survey by Wang et al. [1] reviewed deep learning across fundus photography, OCT, OCTA, and clinical metadata, identifying cross-modal feature fusion as an open research direction with clinical impact. This motivates our exploration of multimodal fusion for cataract segmentation specifically.

### C. Large Pre-Trained Models for Ophthalmic Segmentation

Recent advances in large pre-trained models have changed medical image segmentation. Ma et al. introduced MedSAM [8], fine-tuning the Segment Anything Model (SAM) on over 1.5 million medical image-mask pairs across 10 imaging modalities. While MedSAM achieves strong generalization, its scale, requiring 20 A100 GPUs for training, makes it impractical for deployment on small ophthalmic datasets without substantial computational infrastructure.

The challenge of adapting large models to specific tasks was addressed by Chen et al. with SAM-Adapter [9], which introduces lightweight task-specific adapters to guide SAM toward domain-specific features. While adapter-based methods reduce fine-tuning cost, they still inherit the full SAM encoder and are optimized for scenarios with large pre-training data available.

### D. Positioning of the Proposed Work

The above literature reveals three complementary trends: (1) task-specific architectures like U-Net achieve strong segmentation performance but only have access to the information in a single image modality; (2) cross-modal attention mechanisms can effectively fuse different types of information but have not been applied to combining an image with its own edge map for cataract segmentation; and (3) large pre-trained models provide broad generalization but require computational resources that are impractical for small, specialized datasets.

Our proposed architecture bridges the gap between (1) and (2): we adapt the dual-encoder, cross-attention paradigm demonstrated in mmFormer and DECTNet to a lightweight model where the two encoders process the original RGB image and its Canny edge map respectively. This design explicitly provides the decoder with complementary color and geometric features through a bottleneck cross-attention module, without requiring large-scale pre-training. Compared to large pre-trained model adaptation (3), this approach is computationally accessible and directly optimized for the structural characteristics of cataract anterior segment images.

## III. PROPOSED METHOD

### A. Overview

We propose a dual-encoder segmentation architecture that fuses two complementary representations of the same input image: the original RGB image and a structural edge map derived from it via Canny filtering. Although both inputs originate from a single image, they occupy different feature spaces, color and texture versus geometry, and are processed by independent encoders whose representations are fused through cross-attention modules at multiple scales.

### B. Input Representations

Given an input image  $\mathbf{I} \in \mathbb{R}^{H \times W \times 3}$ , two representations are constructed:

**RGB stream.** The original image is normalized to  $[0, 1]$  and resized to  $512 \times 512$ .




Fig. 1. Proposed dual-encoder architecture. Blue: RGB encoder (Enc-A). Red: Edge encoder (Enc-B). Orange: Cross-attention fusion modules with learnable gates  $\alpha_1, \alpha_2$ . Purple: fused skip projections ( $1 \times 1$  conv). Green: decoder. Dashed arrows indicate edge skip connections concatenated with RGB skips before projection.

**Edge stream.** A single-channel edge map  $\mathbf{E} \in \mathbb{R}^{H \times W \times 1}$  is computed by applying Canny edge detection to the grayscale conversion of  $\mathbf{I}$ . The Canny algorithm works in four steps: (1) smooth the image with a Gaussian filter to remove noise, (2) find the intensity gradient at every pixel (how fast and in what direction brightness changes), (3) thin the resulting edges by keeping only pixels that are local maxima in the gradient direction, and (4) link edge pixels into continuous contours using two thresholds, pixels above  $t_2 = 150$  are kept as strong edges, pixels below  $t_1 = 50$  are discarded, and pixels in between are kept only if they connect to a strong edge. This dual-threshold strategy produces clean, connected boundaries while suppressing noise. After detection, the edge map is replicated to three channels to match the expected input shape of the encoder. This representation makes the geometric boundary structure explicit, which is clinically relevant for lens delineation under variable illumination.




Fig. 2. Sample anterior segment images (top) and their corresponding Canny edge maps (bottom,  $t_1=50, t_2=150$ ). The edge stream makes lens and corneal boundaries geometrically explicit even when color contrast is low.

### C. Dual-Encoder Architecture

Both streams are processed by independent U-Net encoders sharing the same ResNet-34 architecture but no weights.

The encoder follows the standard downsampling path of five resolution stages, producing feature maps at channel depths 64, 64, 128, 256, and 512 respectively. For a  $512 \times 512$  input, the bottleneck feature map has spatial resolution  $16 \times 16$ .

Let  $\mathbf{F}_A^{(s)}$  and  $\mathbf{F}_B^{(s)}$  denote the feature maps from the RGB and edge encoders at stage  $s$ , with the bottleneck at stage 5 ( $C = 512$ ) and stage 4 at ( $C = 256$ , spatial  $32 \times 32$ ).

### D. Cross-Attention Fusion Module

The bottleneck features are fused via multi-head cross-attention [7]. Both feature maps are reshaped from  $(B, C, H', W')$  to sequence format  $(B, H'W', C)$ . The RGB features serve as queries and the edge features supply keys and values:

$$\tilde{\mathbf{F}}_A = \text{LayerNorm}(\alpha \cdot \text{MHA}(Q=\mathbf{F}_A, K=\mathbf{F}_B, V=\mathbf{F}_B) + \mathbf{F}_A) \quad (1)$$

Intuitively, this attention operation lets each spatial location in the RGB feature map “look at” all spatial locations in the edge feature map and decide which boundary cues are relevant at that position. The queries  $\mathbf{Q}$  from the RGB features represent “what boundary information do I need here?” while the keys  $\mathbf{K}$  from the edge features represent “what boundaries are present at each location?” The dot product  $\mathbf{Q}\mathbf{K}^T$  produces an attention map that determines how much each edge location should influence each RGB location.

where  $\alpha = \sigma(\theta)$  is a learnable scalar gate with  $\theta$  initialized to  $-4$ , giving  $\alpha \approx 0.02$  at the start of training. This allows the model to begin as a near-pure RGB U-Net and gradually learn to incorporate edge information only when it is beneficial, an important safeguard given the noisy nature of Canny edge maps. Intuitively, the gate acts as a volume knob for edge information: it starts at zero so the model is not forced to use potentially noisy edge features, and training can turn it up only if the edge cues consistently improve the segmentation.

### E. Multi-Scale Fusion

Fusing edge information only at the bottleneck (the lowest, most compressed resolution) means the decoder receives boundary cues only after they have been heavily downsampled. Fine-grained edge details lost during downsampling cannot be recovered. We address this with two additional mechanisms that deliver edge information at multiple resolutions:

**Stage-4 cross-attention.** A second cross-attention module (256 channels,  $n_{\text{heads}} = 8$ , with its own learnable gate) is applied one encoder stage above the bottleneck ( $32 \times 32$  spatial resolution). This injects geometric cues into the feature hierarchy before the bottleneck, providing richer structural context at an earlier stage of the downsampling path.

**Edge encoder skip connections.** At each of the four skip-connection levels in the decoder, the corresponding edge encoder feature map  $\mathbf{F}_B^{(s)}$  is concatenated with the RGB skip  $\mathbf{F}_A^{(s)}$  along the channel dimension, then projected back to the original channel count via a  $1 \times 1$  convolution:

$$\hat{\mathbf{F}}^{(s)} = \text{Conv}_{1 \times 1} \left( \left[ \mathbf{F}_A^{(s)} \parallel \mathbf{F}_B^{(s)} \right] \right), \quad s \in \{1, 2, 3, 4\} \quad (2)$$

This propagates geometric boundary cues through the full decoder path rather than only at the bottleneck, enabling boundary-aware feature refinement at every decoder resolution.

#### F. Decoder and Output

The decoder follows the standard U-Net expanding path with bilinear upsampling and two  $3 \times 3$  convolutional blocks at each stage. The fused skip connections  $\hat{\mathbf{F}}^{(s)}$  replace the plain RGB skips at each resolution level. A final  $1 \times 1$  convolution with sigmoid activation produces the binary segmentation mask  $\mathbf{M} \in [0, 1]^{H \times W}$ .

#### G. Training Objective

The model is trained with a combined loss:

$$\mathcal{L} = \mathcal{L}_{\text{BCE}} + \mathcal{L}_{\text{Dice}} \quad (3)$$

where  $\mathcal{L}_{\text{BCE}}$  is binary cross-entropy and  $\mathcal{L}_{\text{Dice}} = 1 - \frac{2 \sum \hat{M} \cdot M}{\sum \hat{M} + \sum M}$ . This combination penalizes both pixel-level errors and overlap errors, which is standard practice in medical image segmentation.

### IV. EXPERIMENTAL DESIGN

#### A. Dataset

We evaluate on the Cataract-Seg dataset [4], a publicly available collection of anterior segment images annotated for pixel-wise segmentation of cataract-related structures including the cornea, pupil, and lens. The dataset provides pixel-level masks suitable for binary and multi-class segmentation evaluation. This domain is closely related to the surgical scene annotations established in CaDIS [6], but focuses on still clinical images rather than intraoperative video frames.

The dataset is partitioned into training, validation, and test sets following a 70%/15%/15% stratified split, yielding 210 training, 45 validation, and 45 test images. No patient-level overlap exists across splits.

#### B. Preprocessing

All images are resized to  $512 \times 512$  pixels and RGB values normalized to  $[0, 1]$ . Edge maps are computed per image using Canny filtering applied to the grayscale conversion, with thresholds  $t_1 = 50$  and  $t_2 = 150$ , and replicated to three channels before input to Encoder B. This preprocessing is applied at both training and test time. Preprocessing is applied identically across all baseline and proposed models.

#### C. Data Augmentation

During training, the following augmentations are applied on-the-fly to both the RGB image and its corresponding mask: random horizontal flip ( $p = 0.5$ ), random rotation within  $\pm 15^\circ$ , and random brightness and contrast jitter (factor range  $[0.8, 1.2]$ ). Edge maps are recomputed from the augmented RGB image after geometric transforms. No augmentation is applied at validation or test time.

TABLE I  
MODELS COMPARED IN THE EVALUATION.

| Model                | Description                                    |
|----------------------|------------------------------------------------|
| U-Net (RGB)          | Standard U-Net, RGB input only                 |
| U-Net (Edge)         | Standard U-Net, edge map input only            |
| U-Net (Early Fusion) | Single encoder, RGB+Edge concatenated at input |
| Proposed             | Dual encoder + bottleneck cross-attention      |

#### D. Baselines

Four models are evaluated to isolate the contribution of cross-attention fusion:

All models use the same U-Net backbone, optimizer, and training schedule to ensure fair comparison. The first three baselines share the proposed model's decoder and loss function.

#### E. Metrics

Model performance is reported using Intersection over Union (IoU), Dice coefficient, and F1 score. For binary segmentation, let  $\hat{M}$  be the predicted mask and  $M$  the ground-truth mask. IoU measures the overlap between prediction and ground truth divided by their total area:

$$\text{IoU} = \frac{|\hat{M} \cap M|}{|\hat{M} \cup M|} = \frac{\text{overlap}}{\text{union}} \quad (4)$$

A value of 1 means perfect overlap, 0 means no overlap. The Dice coefficient is a related measure that weights overlap more heavily:

$$\text{Dice} = \frac{2|\hat{M} \cap M|}{|\hat{M}| + |M|} \quad (5)$$

Both metrics penalize false positives (predicting a pixel that is not in the ground truth) and false negatives (missing a pixel that is in the ground truth). IoU is the primary metric, as it is more conservative ( $\text{IoU} \leq \text{Dice}$ ) and standard in segmentation benchmarks. F1 is mathematically equivalent to Dice in binary segmentation and reported in the same column.

#### F. Implementation Details

Models are implemented in PyTorch using `segmentation-models-pytorch` for the U-Net backbone. Cross-attention is implemented with `torch.nn.MultiheadAttention` ( $n_{\text{heads}} = 8$ ). All models are trained with AdamW ( $\text{lr} = 10^{-4}$ , weight decay  $= 10^{-2}$ ) for 100 epochs with a batch size of 8. The combined BCE+Dice loss is used throughout. Metrics are computed with `torchmetrics`. Experiments were run on a single NVIDIA A100 GPU (40 GB) via Google Colab. The random seed is fixed to 42 for all runs; code is available at <https://github.com/DaniloDuque/multimodal-cataract-seg>.

### V. RESULTS

#### A. Quantitative Evaluation

Table II reports IoU, Dice coefficient, and F1 score for all four models on the held-out test set. The proposed model achieves an IoU of 0.9530, matching the best-performing

baseline (U-Net Early Fusion, 0.9532) and outperforming the standard RGB U-Net baseline (0.9529). The edge-only model performs substantially worse (IoU 0.8817), confirming that Canny edge maps do not carry enough color and texture information by themselves but do carry complementary structural cues that benefit the proposed fusion approach.

TABLE II

TEST SET PERFORMANCE OF ALL MODELS. IOU IS THE PRIMARY METRIC. BEST RESULT PER COLUMN IN **BOLD**. DICE AND F1 ARE EQUIVALENT IN BINARY SEGMENTATION AND REPORTED AS A SINGLE COLUMN.

| Model                | IoU           | Dice / F1     |
|----------------------|---------------|---------------|
| U-Net (RGB)          | 0.9529        | 0.9759        |
| U-Net (Edge)         | 0.8817        | 0.9371        |
| U-Net (Early Fusion) | <b>0.9532</b> | <b>0.9761</b> |
| Proposed             | 0.9530        | 0.9759        |

Fig. 3 illustrates the IoU gap between the edge-only model and the three RGB-based models, highlighting that the top three models converge to near-identical performance on this dataset.




Fig. 3. Test set IoU for all four models. Y-axis starts at 0.85 to magnify differences. The proposed model matches the best baseline within 0.0002 IoU.

### B. Training Dynamics

Fig. 4 shows the training and validation loss curves (BCE + Dice) over 100 epochs. All models converge smoothly without divergence. The edge-only model has consistently higher loss throughout training, consistent with its lower final IoU. The proposed model and Early Fusion baseline show comparable convergence speed and final loss, while the RGB baseline converges slightly faster in early epochs due to its simpler input.




Fig. 4. Training (left) and validation (right) loss curves for all four models over 100 epochs.

Fig. 5 shows the validation IoU over training epochs. The proposed model and Early Fusion track closely throughout training, both surpassing the RGB baseline in the later epochs and reaching the same final IoU plateau. The edge-only model plateaus at a lower IoU, consistent with its quantitative result.



Fig. 5. Validation IoU over 100 training epochs for all four models.

### C. Qualitative Analysis

Fig. 6 shows predicted segmentation masks from all four models on three representative test images. The edge-only model produces visibly coarser and noisier boundaries. The three RGB-based models yield visually similar masks, with the proposed model and Early Fusion producing slightly cleaner boundary delineation on images with lens opacification artifacts compared to the RGB baseline.




Fig. 6. Qualitative comparison of predicted masks on three test images. Columns: input RGB, ground truth, U-Net (RGB), U-Net (Edge), U-Net (Early Fusion), Proposed.

### D. Discussion

The near-identical performance of the top three models (IoU range 0.9529–0.9532) reflects dataset saturation rather than architectural differences. With only 30 test images, a difference of 0.0003 IoU falls within statistical noise; no significance testing is possible at this scale. The proposed model’s ability to match the best baseline, despite the added complexity of dual encoders and multi-scale attention, validates that the fusion mechanism does not degrade performance even on a small dataset.

The learnable gate in the cross-attention module is key to this robustness: by initializing the edge contribution near-zero, the model can selectively learn to rely on structural cues only

where they are informative, effectively falling back to RGB-only behavior when edge features are noisy. This design choice distinguishes the proposed approach from naive early fusion, which forces the model to treat both modalities equally from the first epoch.

## VI. CONCLUSION

We presented a dual-encoder segmentation architecture that fuses RGB color features with Canny-derived geometric edge features through multi-scale cross-attention. Three design choices distinguish the proposed model from a naive bottleneck fusion baseline: a learnable scalar gate that allows the model to down-weight noisy edge contributions during training, a second cross-attention module at encoder stage 4 for earlier structural injection, and edge encoder skip connections that propagate boundary cues through the full decoder path.

On the Cataract-Seg dataset, the proposed model achieves an IoU of 0.9530, matching the best-performing baseline (Early Fusion, 0.9532). The results demonstrate that multi-scale cross-attention fusion is competitive with simpler fusion strategies, and that the gated residual formulation prevents performance degradation when the second modality is noisy, a practical advantage in clinical settings where edge map quality may vary with imaging conditions.

The primary limitation of this evaluation is dataset size: with 30 test images, the differences among the top three models are within statistical noise. Future work should validate the proposed architecture on larger ophthalmic datasets and explore learnable edge detectors (e.g., Sobel filters with trainable weights) as a replacement for fixed Canny filtering, which may allow the edge stream to capture more task-relevant structural features. Extension to multi-class segmentation of individual anterior segment structures (cornea, pupil, lens separately) and to intraoperative video frames [6] are also natural next steps.

## REFERENCES

- [1] S. Wang, X. He, Z. Jian, J. Li, C. Xu, Y. Chen, Y. Liu, H. Chen, C. Huang, J. Hu, and Z. Liu, "Advances and prospects of multi-modal ophthalmic artificial intelligence based on deep learning: a review," *Eye and Vision*, vol. 11, no. 1, p. 38, 2024.
- [2] Y. Zhang, N. He, J. Yang, Y. Li, D. Wei, Y. Huang, Y. Zhang, Z. He, and Y. Zheng, "mmFormer: Multimodal medical transformer for incomplete multimodal learning of brain tumor segmentation," in *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, 2022.
- [3] B. Li, Y. Xu, Y. Wang, and B. Zhang, "DECTNet: Dual encoder network combined convolution and transformer architecture for medical image segmentation," *PLOS ONE*, vol. 19, no. 4, p. e0301019, 2024.
- [4] M. Risma, "Cataract-seg dataset," Roboflow Universe, 2023. [Online]. Available: <https://universe.roboflow.com/muhammad-risma/cataract-seg/dataset/1>
- [5] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional networks for biomedical image segmentation," in *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, ser. Lecture Notes in Computer Science, vol. 9351. Springer, 2015, pp. 234–241.
- [6] M. Grammatikopoulou *et al.*, "CaDIS: Cataract dataset for surgical RGB-image segmentation," *Medical Image Analysis*, 2021.
- [7] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, "Attention is all you need," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017.
- [8] J. Ma, Y. He, F. Li, L. Han, C. You, and B. Wang, "Segment anything in medical images," *Nature Communications*, vol. 15, p. 654, 2024.

- [9] T. Chen, L. Zhu, C. Ding, R. Cao, Y. Zhang, Z. Li, L. Kong, P. Jin, and L. Ma, "SAM-Adapter: Adapting segment anything in underperformed scenes," in *Proceedings of the IEEE/CVF International Conference on Computer Vision Workshops (ICCVW)*, 2023.